

CMSC 733: Assignment #2

Due on Thursday, October 15, 2015

Aloimonos, Yiannis 11:00am

Kanishka Ganguly

10/15/2015

Contents

Problem 1	3
(a)	3
(b)	5
(c)	7
(d)	9
Problem 2	12
(a)	12
(b)	13
(c)	13
(d)	14
(e)	14
(f)	14
Problem 3	15
Problem 4	16
(a)	16
(b)	16
(c)	16
Problem 5	17
Problem 6	17
Problem 7	21
Problem 8	21
Problem 9	26
Problem 10	27
(a)	27
(b)	27
(c)	27
(d)	27
Problem 11	28

Problem 1

(a)



Rotation of Image

Listing 1: Rotation

```
1 clear all
2 close all
3 clc
4
5 img = imread('lena.jpg');
6 figure
7 imshow(img);
8 title('Original Image');
9
10 [rows, cols] = size(img);
11
12 center_x = rows/2;
13 center_y = cols/2;
14
15 % Rotation Nearest Neighbour %
16 new_img = zeros(rows);
17 angle = 15;
18 for i = 1:rows
19     for j = 1:cols
20         x = cosd(angle) * (i-center_x) - sind(angle)*(j-center_y) + center_x;
21         y = sind(angle) * (i-center_x) + cosd(angle)*(j-center_y) + center_y;
22         x_2 = round(x);
23         y_2 = round(y);
```


```

24
25     if (x_2 > cols || x_2 < 1)
26         continue
27     elseif(y_2 > rows || y_2 < 1)
28         continue
29     else
30         new_img(i,j) = img(x_2,y_2);
31     end
32 end
33 end
34 rot_img = uint8(new_img);
35
36 % Rotation Bilinear %
37 angle = 15;
38 new_img2 = zeros(rows);
39 for i = 1:rows
40     for j = 1:cols
41         x = cosd(angle) * (i-center_x) - sind(angle)*(j-center_y) + center_x;
42         y = sind(angle) * (i-center_x) + cosd(angle)*(j-center_y) + center_y;
43
44         x1 = floor(x);
45         y1 = floor(y);
46         x2 = ceil(x);
47         y2 = ceil(y);
48
49         if(x1 > cols || x1 < 1)
50             continue
51         elseif(x2 > cols || x2 < 1)
52             continue
53         elseif(y1 > cols || y1 < 1)
54             continue
55         elseif(y2 > cols || y2 < 1)
56             continue
57         else
58             I1 = (x-x1)*img(x2,y1);
59             I2 = (x2-x)*img(x1,y1);
60             I1_ = I1 + I2;
61
62             I1 = (x-x1)*img(x2,y2);
63             I2 = (x2-x)*img(x1,y2);
64             I2_ = I1 + I2;
65
66             I = (y-y1)*I2_ + (y2-y)*I1_;
67             new_img2(i,j) = I;
68         end
69     end
70 end
71 rot_img2 = uint8(new_img2);
72
73 figure
74 subplot(1,2,1)
75 imshow(rot_img);
76 title('Nearest Neighbour Rotation');

```

```
77 subplot(1,2,2)
78 imshow(rot_img2);
79 title('Bilinear Rotation');
```

(b)



Nearest Neighbour Scaling of Image



Bilinear Scaling of Image

Listing 2: Scaling

```
1 clear all
2 close all
3 clc
4
5 img = imread('lena.jpg');
6
7 [rows, cols] = size(img);
8
9 scale = 3;
10 scale_mat = ([scale 0; 0 scale]^-1);
11 rows = round(rows * scale);
12 cols = round(cols * scale);
13
14 % Scaling - Nearest Neighbour %
15 new_img = zeros(rows,cols);
16
17 for i = 1:rows
18     for j = 1:cols
19         coord = scale_mat*[i;j];
20         x_2 = round(coord(1));
21         y_2 = round(coord(2));
22
23         if (x_2 > 0 && y_2 > 0)
```

```

24         new_img(i,j) = img(x_2,y_2);
25     end
26 end
27 end
28 scale_img = uint8(new_img);
29
30 % Scaling Bilinear %
31 new_img2 = zeros(rows,cols);
32
33 for i = 1:rows
34     for j = 1:cols
35         coord = scale_mat*[i;j];
36         x = coord(1);
37         y = coord(2);
38
39         x1 = floor(x);
40         y1 = floor(y);
41         x2 = ceil(x);
42         y2 = ceil(y);
43
44         if (x1 > 0 && x2 > 0 && y1 > 0 && y2 > 0)
45             I1 = (x-x1)*img(x2,y1);
46             I2 = (x2-x)*img(x1,y1);
47             I1_ = I1 + I2;
48
49             I1 = (x-x1)*img(x2,y2);
50             I2 = (x2-x)*img(x1,y2);
51             I2_ = I1 + I2;
52
53             I = (y-y1)*I2_ + (y2-y)*I1_;
54             new_img2(i,j) = I;
55         end
56     end
57 end
58 scale_img2 = uint8(new_img2);
59
60 figure
61 imshow(scale_img);
62 title('Nearest Neighbour Scaling by 3');
63
64 figure
65 imshow(scale_img2);
66 title('Bilinear Scaling by 3');

```

(c)

Listing 3: Skewing

```

1 clear all
2 close all
3 clc
4
5 img = imread('lena.jpg');

```

```
6 figure
7 imshow(img);
8 title('Original Image');
9
10 [rows, cols] = size(img);
11
12 skew = 5;
13
14 % Skewing Nearest Neighbour %
15 new_img = zeros(rows, cols);
16 for i = 1:rows
17     for j = 1:cols
18         x = i + j*tand(skew);
19         y = j;
20         x_2 = round(x);
21         y_2 = round(y);
22
23         if (x_2 > cols || x_2 < 1)
24             continue
25         elseif(y_2 > rows || y_2 < 1)
26             continue
27         else
28             new_img(i, j) = img(x_2, y_2);
29         end
30     end
31 end
32 skew_img = uint8(new_img);
33 figure
34 imshow(skew_img);
35 title('Nearest Neighbour Skewing');
36
37 % Skewing Bilinear %
38 new_img2 = zeros(rows, cols);
39 for i = 1:rows
40     for j = 1:cols
41         x = i + j*tand(skew);
42         y = j;
43
44         x1 = floor(x);
45         y1 = floor(y);
46         x2 = ceil(x);
47         y2 = ceil(y);
48
49         if(x1 > cols || x1 < 1)
50             continue
51         elseif(x2 > cols || x2 < 1)
52             continue
53         elseif(y1 > cols || y1 < 1)
54             continue
55         elseif(y2 > cols || y2 < 1)
56             continue
57         else
58             I1 = (x-x1)*img(x2, y1);
```



```

59         I2 = (x2-x)*img(x1,y1);
60         I1_ = I1 + I2;
61
62         I1 = (x-x1)*img(x2,y2);
63         I2 = (x2-x)*img(x1,y2);
64         I2_ = I1 + I2;
65
66         I = (y-y1)*I2_ + (y2-y)*I1_;
67         new_img2(i,j) = I;
68     end
69 end
70 end
71 skew_img2 = uint8(new_img2);
72 figure
73 imshow(skew_img2);
74 title('Bilinear Skewing');

```

(d)


Listing 4: Affine Transformation

```

1  clear all
2  close all
3  clc
4
5  img = imread('lena.jpg');
6  figure
7  imshow(img);
8  title('Original Image');
9
10 [rows, cols] = size(img);
11
12 T = [1.9 -0.5 0; 0 0 1; 1 1 1] * [1 -1 0; 0 0 3; 1 1 1]^-1
13
14 % Affine - Nearest Neighbour %
15 new_img = zeros(rows,cols);
16
17 for i = 1:rows
18     for j = 1:cols
19         coord = T(1:2,1:2)*[i;j];
20         x_2 = round(coord(1));
21         y_2 = round(coord(2));
22
23         if (x_2 > 0 && x_2 < rows && y_2 > 0 && y_2 < cols)
24             new_img(i,j) = img(x_2,y_2);
25         end
26     end
27 end
28 affine_img = uint8(new_img);
29 figure
30 imshow(affine_img);
31 title('Nearest Neighbour Affine');

```

```
32
33 % Affine - Bilinear %
34 new_img2 = zeros(rows,cols);
35
36 for i = 1:rows
37     for j = 1:cols
38         coord = T(1:2,1:2)*[i;j];
39         x = coord(1);
40         y = coord(2);
41
42         x1 = floor(x);
43         y1 = floor(y);
44         x2 = ceil(x);
45         y2 = ceil(y);
46
47         if (x1 > 0 && x2 > 0 && y1 > 0 && y2 > 0 && x1 < rows && y1 < cols && x2 <
            rows && y2 < cols)
48             I1 = (x-x1)*img(x2,y1);
49             I2 = (x2-x)*img(x1,y1);
50             I1_ = I1 + I2;
51
52             I1 = (x-x1)*img(x2,y2);
53             I2 = (x2-x)*img(x1,y2);
54             I2_ = I1 + I2;
55
56             I = (y-y1)*I2_ + (y2-y)*I1_;
57             new_img2(i,j) = I;
58         end
59     end
60 end
61 affine_img2 = uint8(new_img2);
62 figure
63 imshow(affine_img2);
64 title('Bilinear Affine');
```



Nearest Neighbour Affine Transform of Image



Bilinear Affine Transform of Image

Problem 2

Given Image Patch as

$$\begin{bmatrix} 9 & 10 & 9 & 4 & 3 \\ 5 & 7 & 8 & 9 & 3 \\ 4 & 5 & 6 & 8 & 5 \\ 3 & 4 & 5 & 6 & 8 \\ 2 & 3 & 4 & 5 & 6 \end{bmatrix} \quad (1)$$

(a)

Listing 5: 3x3 Gaussian Filter

```

1 clear all;
2 close all;
3 clc;
4
5 img = [9 10 9 4 3; 5 7 8 9 3; 4 5 6 8 5; 3 4 5 6 8; 2 3 4 5 6]
6
7 % Gaussian Filter %
8 gaussian = (1/16)*[1 2 1; 2 4 2; 1 2 1];
9
10 center_img = img(2:4,2:4);
11
12 for i= 1:3
13     for j = 1:3
14         center_img(i,j) = center_img(i,j) * gaussian(i,j);
15     end
16 end
17 filter_val = 0;
18 filter_val = filter_val + sum(center_img(:,1)) + sum(center_img(:,2)) +
    sum(center_img(:,3));
19 gauss_img = img;
20 gauss_img(3,3) = filter_val;
21 gauss_img

```

Listing 6: 3x3 Box Filter

```

1 clear all;
2 close all;
3 clc;
4
5 img = [9 10 9 4 3; 5 7 8 9 3; 4 5 6 8 5; 3 4 5 6 8; 2 3 4 5 6]
6
7 % Box Filter %
8 center_img = img(2:4,2:4);
9 avg = (sum(center_img(:,1)) + sum(center_img(:,2)) + sum(center_img(:,3)))/9;
10 box_img = img;
11 box_img(3,3) = avg;
12 box_img

```

(b)

Listing 7: Sobel Edge Detector

```

1 clear all;
2 close all;
3 clc;
4
5 img = [9 10 9 4 3; 5 7 8 9 3; 4 5 6 8 5; 3 4 5 6 8; 2 3 4 5 6];
6 sobel_x = (1/8)*[-1 0 1; -2 0 2; -1 0 1];
7 sobel_y = (1/8)*[1 2 1; 0 0 0; -1 -2 -1];
8
9 center_img = img(2:4,2:4);
10
11 for i= 1:3
12     for j = 1:3
13         center_img(i,j) = center_img(i,j) * sobel_x(i,j);
14     end
15 end
16 sobel_x_val = sum(center_img(:,1)) + sum(center_img(:,2)) + sum(center_img(:,3));
17
18 center_img = img(2:4,2:4);
19 for i= 1:3
20     for j = 1:3
21         center_img(i,j) = center_img(i,j) * sobel_y(i,j);
22     end
23 end
24 sobel_y_val = sum(center_img(:,1)) + sum(center_img(:,2)) + sum(center_img(:,3));
25
26 edge_direction = atand(sobel_y_val/sobel_x_val)
27 edge_strength = sqrt(sobel_x_val^2 + sobel_y_val^2)

```

We get the following output:

```

edge_direction = 50.1944
edge_strength = 1.9526

```

(c)

Listing 8: Median Filter

```

1 clear all;
2 close all;
3 clc;
4
5 img = [9 10 9 4 3; 5 7 8 9 3; 4 5 6 8 5; 3 4 5 6 8; 2 3 4 5 6]
6
7 % Median Filter %
8 center_img = img(2:4,2:4);
9 arr = center_img(:);
10 median = median(arr);
11 median_img = img;
12 median_img(3,3) = median;
13 median_img

```

Median filtering has been shown to give best results for salt-and-pepper noise. Median filtering has the added advantage (over mean filters) of removing noise while preserving edges in an image (under certain conditions).

- Since the median is a more robust averaging technique than the mean, a single unrepresentative pixel in a neighbourhood will not affect the median value significantly. Thus, it is robust to the presence of outliers.
- Also, since the median value must actually be the value of one of the pixels in the neighborhood, the median filter does not create new unrealistic pixel values when the filter straddles an edge. For this reason the median filter is much better at preserving sharp edges than the mean filter.

(d)

Why the Gaussian is a good smoothing filter

- A Gaussian filter is a better smoothing filter (as compared to the mean filter) because the Gaussian filter outputs a ‘weighted average’ of each pixel’s neighbourhood, where the average is weighted based on the distance each neighbouring pixel is away from the central pixel. So, pixels directly neighbouring the central pixel is weighted more than the ones farther away from the central pixel. This gives a much ‘gentler’ smoothing and preserves edges better, as compared to a mean filter’s uniformly weighted averaging technique.
- Also, the Gaussian filter has a better *frequency response* than the box (mean) filter. While both filters remove high spatial frequency components from an image, the Gaussian filter shows little to no oscillations in its frequency response as compared to the mean filter, which shows several oscillations.

To perform Gaussian filtering fast

- Gaussian filter has the property that it is separable, i.e. we can express the 2-D convolution as two 1-D convolutions. Thus, two 1-D convolutions can be performed in lesser time than one 2-D convolutions.
- If the filter is large, we can use the fact that a convolution in the spatial domain is equivalent to a multiplication in the frequency (Fourier) domain. This works well for multiple images, since we need to take the Fourier Transform of the Gaussian filter only once. Also, the set of Gaussian functions is closed under Fourier transforms, which means that the Fourier transform of a Gaussian is another Gaussian.

(e)

If the Gaussian kernel is larger than the width of the line, then after smoothing, the width of the line is reduced, i.e. the line becomes ‘thinner’. Also, at the edges of the line, we can see a ‘tapering’ effect.

(f)

A box filter attenuates the noise because on taking the mean value of neighbouring pixels, the resulting output reduces the ‘variance’ of the noise.

That is, the variance of noise in the mean is smaller than variance of the pixel noise.

Problem 3

We have a 1-D image as follows:

$$I(x) = \cos(x/100) \tag{2}$$

When we apply gradient smoothing to $I(x)$, there is no effect on the intensity, which remains unchanged. Thus, $I_{smooth} = \cos(x/100)$.

On this smoothed image, when we calculate the derivatives, we use the mask $[-101]$. We see that the maximum absolute value (magnitude) of the derivative is obtained at $x = 100[\pi/2 + n\pi]$. This is because the derivative of $\cos(x)$ is $-\sin(x)$, which is maximum at $\pi/2 + n\pi$. We can thus find the edges for the image at the zero-crossings of $\cos(x/100)$.

Now, we see that the effect of increasing the σ of the Gaussian filter doesn't affect the image since the intensities remain unchanged even after smoothing.

Also, an increase in the threshold value doesn't have any effect on the local maxima of the derivative since the maximum absolute value is still obtained at $x = 100[\pi/2 + n\pi]$.

Problem 4

Let the given image I be

$$\begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ 3 & 4 & 5 & 6 & 7 \\ 5 & 6 & 7 & 8 & 9 \\ 7 & 8 & 9 & 10 & 11 \\ 9 & 10 & 11 & 12 & 13 \end{bmatrix} \quad (3)$$

(a)

For the gradient in the x -direction, we design a filter $K = \frac{1}{2}[-1 \ 0 \ 1]$. Upon convolution with $f(x, y)$ we get $\frac{1}{2}[f(x+1, y) - f(x-1, y)]$ which is the gradient along the x -direction.

$$\therefore \text{gradient}_x = \frac{8-6}{2} = 1 \quad (4)$$

Similarly, to compute the gradient along the y -direction, we choose the filter $M = [-1 \ 0 \ 1]^T$.

$$\therefore \text{gradient}_y = \frac{9-5}{2} = 2 \quad (5)$$

The resultant gradient is given by:

$$\text{gradient}_{total} = \sqrt{\text{gradient}_x^2 + \text{gradient}_y^2} = \sqrt{5} \quad (6)$$

To calculate the direction, we have $\tan^{-1}(2)^\circ$ with x -axis or 63.44° .

(b)

It is given that gradient at point $(3, 7)$ is $(3, -2)$ and the intensity at that point is 17.

We know that the gradient is defined as the change of intensity per unit distance, i.e. per pixel.

Thus we have

$$I(3.1, 7.3) = I(3, 7) + G_x(0.1) + G_y(0.3) \quad (7)$$

$$\implies 17 + 0.3 - 0.6 = 16.7 \quad (8)$$

(c)

It is given that $I(x, y) = (x-5)^2 + (y-1)^2$

At (x, y) ,

$$\text{gradient}_x = \frac{1}{2}[I(x+1, y) - I(x-1, y)] = -4 \quad (9)$$

$$\text{gradient}_y = \frac{1}{2}[I(x, y+1) - I(x, y-1)] = 2 \quad (10)$$

$$(11)$$

Thus we have gradient at $(3, 2)$ as $(-4, 2)$.

Problem 5

NOT ATTEMPTED

Problem 6

Listing 9: Main Routine

```

1  clear all;
2  close all;
3  clc;
4
5  myDir = 'hist_img/';
6  ext_img = '*.jpg';
7  a = dir([myDir ext_img]);
8  nfile = max(size(a));
9  for i=1:nfile
10     my_img1(i).img = imread([myDir a(i).name]);
11     my_img2(i).img = imread([myDir a(i).name]);
12 end
13
14 for i=1:6
15     for j=(i+1):6
16         h1 = Q6_myColorHist(my_img1(i).img,0);
17         h2 = Q6_myColorHist(my_img2(j).img,0);
18         ssd = Q6_histDist(h1, h2);
19         ssd_red_table(i,j) = ssd(1,:);
20         ssd_green_table(i,j) = ssd(2,:);
21         ssd_blue_table(i,j) = ssd(3,:);
22         ssd_red_table(j,i) = ssd(1,:);
23         ssd_green_table(j,i) = ssd(2,:);
24         ssd_blue_table(j,i) = ssd(3,:);
25     end
26 end
27
28 % Semantic Similarity In Percentages Between Images %
29 similarity = Q6_similarity(ssd_red_table,ssd_green_table, ssd_blue_table);
30
31 images_list = {'desert1', 'desert2', 'desert3', 'forest1', 'forest2', 'forest3'};
32 array2table(similarity, 'VariableNames', images_list, 'RowNames', images_list)

```

Listing 10: Function *myColorHist(I)*

```

1  function [ hist ] = Q6_myColorHist( img, show_hist )
2  [rows, cols] = size(img);
3
4  img_red = img(:,:,1);
5  img_green = img(:,:,2);
6  img_blue = img(:,:,3);
7
8  bin_red = zeros(1,8);
9  bin_green = zeros(1,8);
10 bin_blue = zeros(1,8);

```

```
11
12 for (i=1:rows)
13     for (j=1:cols/3)
14         if (img_red(i,j) >= 1 && img_red(i,j) <= 32)
15             bin_red(1) = bin_red(1) + 1;
16         elseif (img_red(i,j) >= 33 && img_red(i,j) <= 64)
17             bin_red(2) = bin_red(2) + 1;
18         elseif (img_red(i,j) >= 65 && img_red(i,j) <= 96)
19             bin_red(3) = bin_red(3) + 1;
20         elseif (img_red(i,j) >= 97 && img_red(i,j) <= 128)
21             bin_red(4) = bin_red(4) + 1;
22         elseif (img_red(i,j) >= 129 && img_red(i,j) <= 160)
23             bin_red(5) = bin_red(5) + 1;
24         elseif (img_red(i,j) >= 161 && img_red(i,j) <= 192)
25             bin_red(6) = bin_red(6) + 1;
26         elseif (img_red(i,j) >= 193 && img_red(i,j) <= 224)
27             bin_red(7) = bin_red(7) + 1;
28         elseif (img_red(i,j) >= 225 && img_red(i,j) <= 255)
29             bin_red(8) = bin_red(8) + 1;
30         end
31     end
32 end
33
34 for (i=1:rows)
35     for (j=1:cols/3)
36         if (img_green(i,j) >= 1 && img_green(i,j) <= 32)
37             bin_green(1) = bin_green(1) + 1;
38         elseif (img_green(i,j) >= 33 && img_green(i,j) <= 64)
39             bin_green(2) = bin_green(2) + 1;
40         elseif (img_green(i,j) >= 65 && img_green(i,j) <= 96)
41             bin_green(3) = bin_green(3) + 1;
42         elseif (img_green(i,j) >= 97 && img_green(i,j) <= 128)
43             bin_green(4) = bin_green(4) + 1;
44         elseif (img_green(i,j) >= 129 && img_green(i,j) <= 160)
45             bin_green(5) = bin_green(5) + 1;
46         elseif (img_green(i,j) >= 161 && img_green(i,j) <= 192)
47             bin_green(6) = bin_green(6) + 1;
48         elseif (img_green(i,j) >= 193 && img_green(i,j) <= 224)
49             bin_green(7) = bin_green(7) + 1;
50         elseif (img_green(i,j) >= 225 && img_green(i,j) <= 255)
51             bin_green(8) = bin_green(8) + 1;
52         end
53     end
54 end
55
56 for (i=1:rows)
57     for (j=1:cols/3)
58         if (img_blue(i,j) >= 1 && img_blue(i,j) <= 32)
59             bin_blue(1) = bin_blue(1) + 1;
60         elseif (img_blue(i,j) >= 33 && img_blue(i,j) <= 64)
61             bin_blue(2) = bin_blue(2) + 1;
62         elseif (img_blue(i,j) >= 65 && img_blue(i,j) <= 96)
63             bin_blue(3) = bin_blue(3) + 1;
```

```

64         elseif(img_blue(i,j) >= 97 && img_blue(i,j) <= 128)
65             bin_blue(4) = bin_blue(4) + 1;
66         elseif(img_blue(i,j) >= 129 && img_blue(i,j) <= 160)
67             bin_blue(5) = bin_blue(5) + 1;
68         elseif(img_blue(i,j) >= 161 && img_blue(i,j) <= 192)
69             bin_blue(6) = bin_blue(6) + 1;
70         elseif(img_blue(i,j) >= 193 && img_blue(i,j) <= 224)
71             bin_blue(7) = bin_blue(7) + 1;
72         elseif(img_blue(i,j) >= 225 && img_blue(i,j) <= 255)
73             bin_blue(8) = bin_blue(8) + 1;
74     end
75 end
76 end
77
78 % Normalization Stage %
79 sum_red = sum(bin_red);
80 bin_red = bin_red/sum_red;
81 sum_green = sum(bin_green);
82 bin_green = bin_green/sum_green;
83 sum_blue = sum(bin_blue);
84 bin_blue = bin_blue/sum_blue;
85 hist = [bin_red;bin_green;bin_blue];
86
87 % Histogram Visualization %
88 if show_hist == 1
89     hist_plot=figure('Position', [100, 100, 3000, 3000]);
90     subplot(1,3,1)
91     bar(bin_red)
92     set(gca,'XTickLabel',{'1-32', '33-64', '65-96', '97-128','129-160',
93         '161-192', '193-224', '225-255'},'FontSize',8)
94     title('Red Channel Histogram')
95     xlabel('Intensity')
96     ylabel('Pixel Count')
97
98     subplot(1,3,2)
99     bar(bin_green)
100    set(gca,'XTickLabel',{'1-32', '33-64', '65-96', '97-128','129-160',
101        '161-192', '193-224', '225-255'},'FontSize',8)
102    title('Green Channel Histogram')
103    xlabel('Intensity')
104    ylabel('Pixel Count')
105
106    subplot(1,3,3)
107    bar(bin_blue)
108    set(gca,'XTickLabel',{'1-32', '33-64', '65-96', '97-128','129-160',
109        '161-192', '193-224', '225-255'},'FontSize',8)
110    title('Blue Channel Histogram')
111    xlabel('Intensity')
112    ylabel('Pixel Count')
113 end
114 end

```

Listing 11: Function *histDist(h1, h2)*

```

1 function [ ssd ] = Q6_histDist( h1, h2 )
2 ssd_red = sum((h1(1,:)-h2(1,:)).^2);
3 ssd_green = sum((h1(2,:)-h2(2,:)).^2);
4 ssd_blue = sum((h1(3,:)-h2(3,:)).^2);
5 ssd = [ssd_red;ssd_green;ssd_blue];
6 end

```

For extra credit, I have developed a function that returns a percentage similarity between images that can be used to determine semantic similarity between scenes.

Listing 12: Function *similarity(ssdR, ssdG, ssdB)*

```

1 function [ similarity ] = Q6_similarity( ssd_red, ssd_green, ssd_blue )
2
3 ssd_red_percent = 1 - ssd_red;
4 ssd_green_percent = 1 - ssd_green;
5 ssd_blue_percent = 1 - ssd_blue;
6
7 similarity = bsxfun(@times, ssd_blue_percent, bsxfun(@times, ssd_red_percent,
8               ssd_green_percent));
9 similarity = similarity * 100;
10 end

```

ans =

	desert1	desert2	desert3	forest1	forest2	forest3
desert1	100	66.418	89.814	47.336	71.69	64.536
desert2	66.418	100	56.619	24.911	49.945	46.145
desert3	89.814	56.619	100	52.115	65.898	62.29
forest1	47.336	24.911	52.115	100	57.396	76.786
forest2	71.69	49.945	65.898	57.396	100	71.908
forest3	64.536	46.145	62.29	76.786	71.908	100

Percentage Matrix for Semantic Similarity Determination Between Scenes

From the individual histograms of the Red, Green and Blue channels of two images, it may be determined whether two images are ‘similar’ in which channel. This, combined with a precomputed set of ‘generic’ values for each type of scene can be used to determine and differentiate between two scenes, such as a desert and a forest.

However, a more intelligent and adaptive algorithm than the one presented here is required, since it is evident from the values that in some cases, the presences of certain ‘unnecessary’ features can dominate a histogram output.

For example, in a desert and a forest scene, if the ‘amount’ of blue skies that are present dominate the ‘actual’ desert and forest portions, then there would be a higher similarity being returned than what is expected.

Thus, along with color histograms, it might be necessary and useful to include other classification algorithms to detect ‘features’ such as sand and trees in desert and forest scenes respectively.

Problem 7

NOT ATTEMPTED

Problem 8

Listing 13: Corner Detection

```

1  clear all;
2  close all;
3  clc;
4
5  img = zeros(500, 500);
6  img(100:200,100) = 255;
7  img(100:200,300) = 255;
8  img(100,100:300) = 255;
9  img(200,100:300) = 255;
10 img_show = uint8(img);
11
12 figure
13 subplot(3,2,[1,2])
14 subimage(img_show)
15 set(gca,'XtickLabel',[],'YtickLabel',[]);
16 title('Original Image');
17
18 top_left_mask = [0 0 0 0 0; 0 0 0 0 0; 0 0 1 1 1; 0 0 1 0 0; 0 0 1 0 0];
19 top_right_mask = [0 0 0 0 0; 0 0 0 0 0; 1 1 1 0 0; 0 0 1 0 0; 0 0 1 0 0];
20 bottom_left_mask = [0 0 1 0 0; 0 0 1 0 0; 0 0 1 1 1; 0 0 0 0 0; 0 0 0 0 0];
21 bottom_right_mask = [0 0 1 0 0; 0 0 1 0 0; 1 1 1 0 0; 0 0 0 0 0; 0 0 0 0 0];
22
23 [rows, cols] = size(img);
24
25 baseline = 0;
26
27 for i=1:rows
28     for j=1:cols
29         if (i == 1 && j == 1)
30             edge = bsxfun(@times,img(i:i+4,j:j+4),top_left_mask);
31             baseline = edge(3,3);
32         end
33         if i+4 < rows && j+4 < cols
34             edge = bsxfun(@times,img(i:i+4,j:j+4),top_left_mask);
35             if edge(3,3) == edge(3,4)
36                 if edge(3,4) == edge(3,5)
37                     if edge(3,5) == edge(4,3)
38                         if edge(4,3) == edge(5,3)
39                             if (edge(3,3) ~= baseline)
40                                 img(i+2,j+2) = 0;
41                                 break;
42                             else
43                                 continue;
44                             end
45                         end
46                     end
47                 end
48             end
49         end
50     end
51 end

```

```

46         end
47     end
48 end
49 end
50 end
51 end
52 img_show = uint8(img);
53 subplot(3,2,3)
54 subimage(img_show)
55 set(gca,'XtickLabel',[],'YtickLabel',[]);
56 title('Top Left Corner Detected');
57
58 img = zeros(500, 500);
59 img(100:200,100) = 255;
60 img(100:200,300) = 255;
61 img(100,100:300) = 255;
62 img(200,100:300) = 255;
63
64 for i=1:rows
65     for j=1:cols
66         if(i == 1 && j == 1)
67             edge = bsxfun(@times,img(i:i+4,j:j+4),top_right_mask);
68             baseline = edge(3,3);
69         end
70         if i+4 < rows && j+4 < cols
71             edge = bsxfun(@times,img(i:i+4,j:j+4),top_right_mask);
72             if edge(3,3) == edge(3,2)
73                 if edge(3,2) == edge(3,1)
74                     if edge(3,1) == edge(4,3)
75                         if edge(4,3) == edge(5,3)
76                             if (edge(3,3) ~= baseline)
77                                 img(i+2,j+2) = 0;
78                                 break;
79                             else
80                                 continue;
81                             end
82                         end
83                     end
84                 end
85             end
86         end
87     end
88 end
89 img_show = uint8(img);
90 subplot(3,2,4)
91 subimage(img_show)
92 set(gca,'XtickLabel',[],'YtickLabel',[]);
93 title('Top Right Corner Detected');
94
95 img = zeros(500, 500);
96 img(100:200,100) = 255;
97 img(100:200,300) = 255;
98 img(100,100:300) = 255;

```

```


99  img(200,100:300) = 255;
100
101  for i=1:rows
102      for j=1:cols
103          if (i == 1 && j == 1)
104              edge = bsxfun(@times,img(i:i+4,j:j+4),top_right_mask);
105              baseline = edge(3,3);
106          end
107          if i+4 < rows && j+4 < cols
108              edge = bsxfun(@times,img(i:i+4,j:j+4),bottom_left_mask);
109              if edge(3,3) == edge(2,3)
110                  if edge(2,3) == edge(1,3)
111                      if edge(1,3) == edge(3,4)
112                          if edge(3,4) == edge(3,5)
113                              if (edge(3,3) ~= baseline)
114                                  img(i+2,j+2) = 0;
115                                  break;
116                              else
117                                  continue;
118                              end
119                          end
120                      end
121                  end
122              end
123          end
124      end
125  end
126  img_show = uint8(img);
127  subplot(3,2,5)
128  subimage(img_show)
129  set(gca,'XtickLabel',[],'YtickLabel',[]);
130  title('Bottom Left Corner Detected');
131  figure
132  imshow(img_show)
133
134  img = zeros(500, 500);
135  img(100:200,100) = 255;
136  img(100:200,300) = 255;
137  img(100,100:300) = 255;
138  img(200,100:300) = 255;
139
140  for i=1:rows
141      for j=1:cols
142          if (i == 1 && j == 1)
143              edge = bsxfun(@times,img(i:i+4,j:j+4),bottom_right_mask);
144              baseline = edge(3,3);
145          end
146          if i+4 < rows && j+4 < cols
147              edge = bsxfun(@times,img(i:i+4,j:j+4),bottom_right_mask);
148              if edge(3,3) == edge(3,2)
149                  if edge(3,2) == edge(3,1)
150                      if edge(3,1) == edge(2,3)
151                          if edge(2,3) == edge(1,3)

```

```

152         if (edge(3,3) ~= baseline)
153             img(i+2,j+2) = 0;
154             break;
155         else
156             continue;
157         end
158     end
159 end
160 end
161 end
162 end
163 end
164 end
165 img_show = uint8(img);
166 subplot(3,2,6)
167 subimage(img_show)
168 set(gca,'XtickLabel',[],'YtickLabel',[]);
169 title('Bottom Right Corner Detected');
170 figure
171 imshow(img_show)

```



All four corners detected and respective pixel removed for visualization

The algorithm proposed and demonstrated above uses the following masks:

$$\text{Top Left Detector} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{bmatrix} \quad (12)$$

$$\text{Top Right Detector} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{bmatrix} \quad (13)$$

$$\text{Bottom Left Detector} = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \quad (14)$$

$$\text{Bottom Right Detector} = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \quad (15)$$

$$(16)$$

The algorithm works as follows:

1. Pad the image to make **ROWS** and **COLS** multiples of 5
2. Place respective mask over image, starting from top-left corner
 - (a) Multiply mask element with image patch
 - (b) Calculate value for element (3,3) and store as **BASELINE**
3. Traverse mask over image from **ROWS** to **COLS**
 - (a) Multiply mask element with image patch
 - (b) Check if all non-zero elements of matrix are equal.
 - (c) IF EQUAL, compare value of non-zero element with **BASELINE**.
 - (d) IF EQUAL, corner is found. Break loop.
4. ELSE, REPEAT.


This algorithm makes the assumption that the rectangle is made of lines which are 1 pixel thick and that the background is of uniform intensity.

It also assumes that the rectangle is anywhere within the bounds of the image and not touching any of the boundaries of the image.

This is a fairly robust algorithm since it can detect corners of multiple rectangles in the same image. Also, the entire algorithm can detect all four corners of a rectangle in one iteration, by placing all four masks in the same loop routine.

The main drawback of said algorithm is with boundary cases where any one edge of the rectangle is touching the edge of the image.

Problem 9



Figures 1. $h(x)$ 2. $f(x)$ 3. $f(x) * h(x)$ 4. $h(x) * h(x)$

Problem 10

(a)

Most images taken in a real-world situation (using normal cameras) have an inherent noise. Also, these images have several features that contain edges.

Taking derivatives using finite differences is an operation that reduces the noise in the image. Also, using a smoothing filter such as a Gaussian causes these edges to blur, thus solving the problem of inhibiting the local edges that are unimportant to the analysis of the image, while retaining the more prominent, useful edges.

(b)

In the former implementation, i.e. convolving with Gaussian first and then taking the Laplacian, we first apply the Gaussian filter to the image using a discrete convolution and then apply the Laplacian operator. This requires two computations to the image, which decreases efficiency since the Gaussian filter needs to be applied to each image that we may have.

However, in the latter implementation, i.e. Laplacian of Gaussian, we can precompute the effect of the Laplacian on the Gaussian filter and then store and apply the resulting filter to the image. This reduces the computation time to just one major calculation, irrespective of the number of images we may have.

(c)

We can show that the Laplacian of the Gaussian is $G_1 * I - G_2 * I$ by taking

$$LoG(x, y) = -\frac{1}{\pi\sigma^4} \left[1 - \frac{x^2 + y^2}{2\sigma^2} \right] e^{-\frac{x^2 + y^2}{2\sigma^2}} \quad (17)$$

This equation can be approximated by using the difference of two Gaussian kernels with suitable values for σ . This ratio turns out to be 1 : 1.6 for the best results.

$$DoG = G_{\sigma_1} - G_{\sigma_2} \quad (18)$$

$$DoG = \frac{1}{2\pi} \left(\frac{1}{\sigma_1} e^{-\frac{x^2 + y^2}{2\sigma_1^2}} - \frac{1}{\sigma_2} e^{-\frac{x^2 + y^2}{2\sigma_2^2}} \right) \quad (19)$$

(d)

Difference of Gaussians (DoG) is more efficient than Laplacian of Gaussian (LoG) because one 2-D Gaussian operation is separable into two 1-D Gaussian operations, which is much more computationally efficient. Also, when transformed to the frequency domain, a Gaussian convolution operation becomes a multiplication operation. This also makes it more computationally efficient to calculate (DoG) than (LoG).

Problem 11

Listing 14: Unsharp Sharpening

```
1 clear all;
2 close all;
3 clc;
4
5 img = imread('lena.jpg');
6 scaling_constant = 0.7;
7
8 % Standard Kernel Size %
9 box = imboxfilt(img,3);
10 gauss = imgaussfilt(img,2);
11
12 img_box_subtract = img - box;
13 img_gauss_subtract = img - gauss;
14
15 img_box_sharpen = img + (scaling_constant * img_box_subtract);
16 img_gauss_sharpen = img + (scaling_constant * img_gauss_subtract);
17
18 figure
19 subplot(1,2,1)
20 imshow(img)
21 title('Original Image')
22 subplot(1,2,2)
23 imshow(img_box_sharpen)
24 title('Box Sharpened Image')
25
26 figure
27 subplot(1,2,1)
28 imshow(img)
29 title('Original Image')
30 subplot(1,2,2)
31 imshow(img_gauss_sharpen)
32 title('Gaussian Sharpened Image')
33
34 clear all;
35 close all;
36 clc;
37 img = imread('lena.jpg');
38 scaling_constant = 0.7;
39
40 % Increased Kernel Size %
41 box = imboxfilt(img,7);
42 gauss = imgaussfilt(img,4);
43
44 img_box_subtract = img - box;
45 img_gauss_subtract = img - gauss;
46
47 img_box_sharpen = img + (scaling_constant * img_box_subtract);
48 img_gauss_sharpen = img + (scaling_constant * img_gauss_subtract);
49
50 figure
```

```
51 subplot(1,2,1)
52 imshow(img)
53 title('Original Image')
54 subplot(1,2,2)
55 imshow(img_box_sharpen)
56 title('Box Sharpened Image')
57
58 figure
59 subplot(1,2,1)
60 imshow(img)
61 title('Original Image')
62 subplot(1,2,2)
63 imshow(img_gauss_sharpen)
64 title('Gaussian Sharpened Image')
```



Unsharp Sharpening Using 3x3 Box Filter

By increasing kernel size, we get



Unsharp Sharpening Using Gaussian Filter with $\sigma = 4$



Unsharp Sharpening Using 7x7 Box Filter